

```
In [1]: import pandas as pd
```

```
# Load the uploaded dataset
file_path = "Flyzy Flight Cancellation - Sheet1.csv"
df = pd.read_csv(file_path)

# Display the first few rows and summary info
df.head(), df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3000 entries, 0 to 2999
Data columns (total 14 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Flight ID                            3000 non-null   int64
1   Airline                             3000 non-null   object
2   Flight_Distance                      3000 non-null   int64
3   Origin_Airport                      3000 non-null   object
4   Destination_Airport                 3000 non-null   object
5   Scheduled_Departure_Time             3000 non-null   int64
6   Day_of_Week                         3000 non-null   int64
7   Month                               3000 non-null   int64
8   Airplane_Type                       3000 non-null   object
9   Weather_Score                       3000 non-null   float64
10  Previous_Flight_Delay_Minutes        3000 non-null   float64
11  Airline_Rating                       3000 non-null   float64
12  Passenger_Load                       3000 non-null   float64
13  Flight_Cancelled                     3000 non-null   int64
dtypes: float64(4), int64(6), object(4)
memory usage: 328.3+ KB
```

```
Out[1]: (
  Flight ID      Airline  Flight_Distance  Origin_Airport  Destination_Airport  \
0      7319483  Airline D              475          Airport 3          Airport 2
1      4791965  Airline E              538          Airport 5          Airport 4
2      2991718  Airline C              565          Airport 1          Airport 2
3      4220106  Airline E              658          Airport 5          Airport 3
4      2263008  Airline E              566          Airport 2          Airport 2

  Scheduled_Departure_Time  Day_of_Week  Month  Airplane_Type  Weather_Score  \
0                      4              6      1          Type C      0.225122
1                      12              1      6          Type B      0.060346
2                      17              3      9          Type C      0.093920
3                      1              1      8          Type B      0.656750
4                      19              7     12          Type E      0.505211

  Previous_Flight_Delay_Minutes  Airline_Rating  Passenger_Load  \
0                      5.0          2.151974          0.477202
1                      68.0          1.600779          0.159718
2                      18.0          4.406848          0.256803
3                      13.0          0.998757          0.504077
4                      4.0          3.806206          0.019638

  Flight_Cancelled
0                0
1                1
2                0
3                1
4                0 ,
None)
```

```
In [2]: import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from scipy.stats import zscore

# Select numeric columns for outlier detection
numeric_cols = df.select_dtypes(include=[np.number]).columns.tolist()
numeric_cols.remove("Flight ID") # Exclude ID column

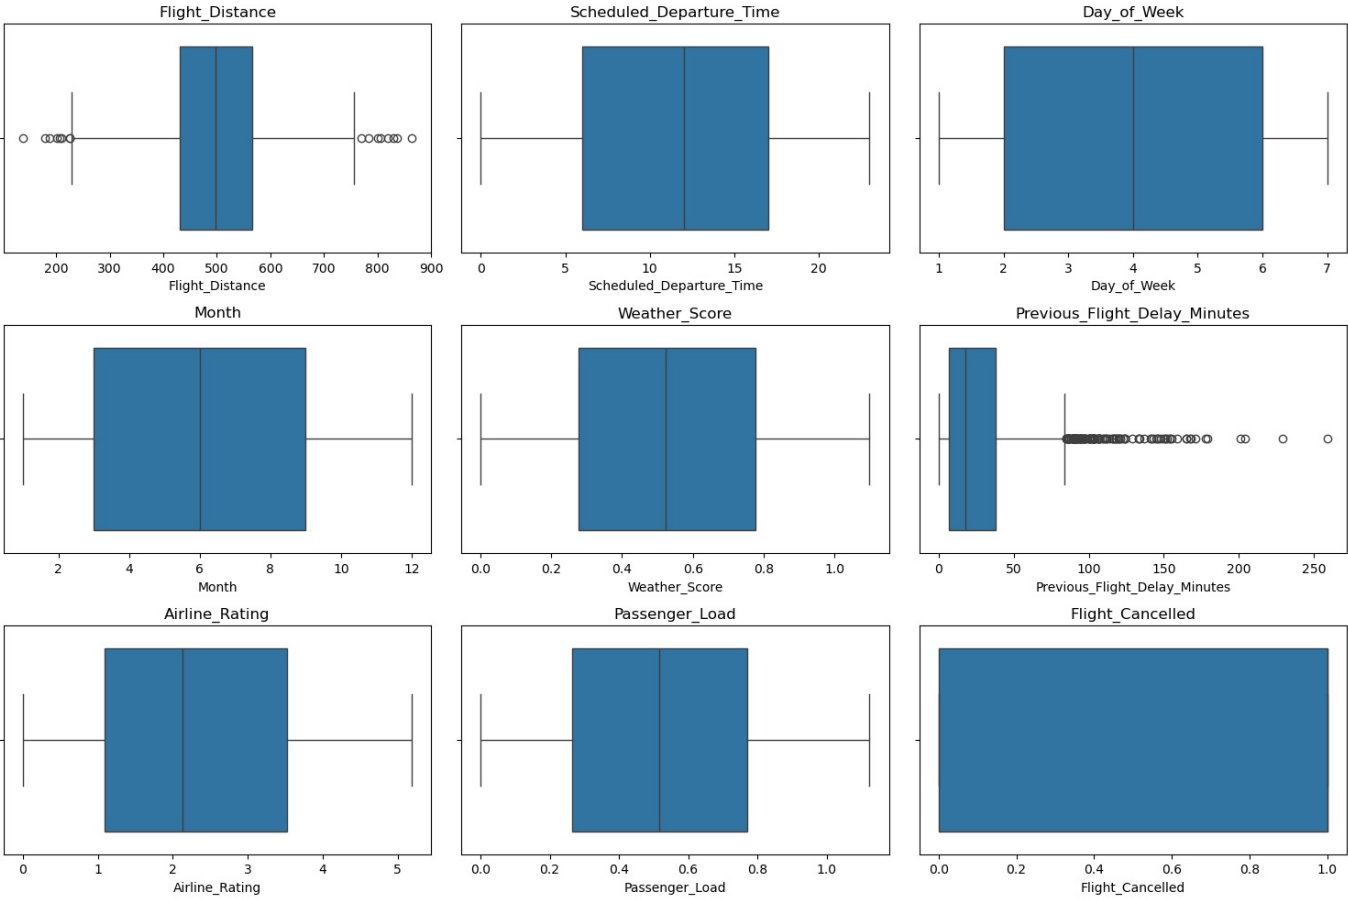
# Compute Z-scores
z_scores = np.abs(zscore(df[numeric_cols]))
outliers = (z_scores > 3)

# Count outliers per column
outlier_counts = outliers.sum(axis=0)

# Visualize outliers with boxplots
plt.figure(figsize=(15, 10))
for i, col in enumerate(numeric_cols, 1):
    plt.subplot(3, 3, i)
    sns.boxplot(x=df[col])
    plt.title(col)
```

```
plt.tight_layout()
plt.show()
```

outlier_counts



```
Out[2]: Flight_Distance      10
Scheduled_Departure_Time    0
Day_of_Week                 0
Month                      0
Weather_Score               0
Previous_Flight_Delay_Minutes 51
Airline_Rating              0
Passenger_Load              0
Flight_Cancelled            0
dtype: int64
```

In []: